



Python's Data Science Stack

Jake VanderPlas @jakevdp
JSM, July 31, 2016

Python is *not* a statistical computing language!

**Python is *not* a statistical
computing language!**


**. . . and this may be its
greatest strength as a
language for statistical
computing.**

A Quick Tour of Python's Data Science Stack

Python's Data Science Stack



IP[y]:
IPython



NumPy



Cython




python™



NumPy = Numerical Python

Efficient array storage, manipulation, and computation



NumPy python

IP[y]:
IPython

python™



NumPy = Numerical Python

Efficient array storage, manipulation, and computation

```
In [1]: import numpy as np
```

```
# Create a 5x5 uniform random matrix
```


```
M = np.random.rand(5, 5)
```

```
# Compute the SVD
```

```
U, S, VT = np.linalg.svd(M)
```

```
print(S)
```

```
[ 2.46102945  0.94542853  0.53550015  0.20705388  0.13071452]
```



NumPy



IP[y]:
IPython

python™






IP[y]:
IPython



Cython = C + Python

Super-set of the Python language that allows easy interfacing with C & Fortran libraries (e.g. BLAS, LAPACK, etc.) and also fast Python code.

Drives many of the packages in the data science stack.






IP[y]:
IPython

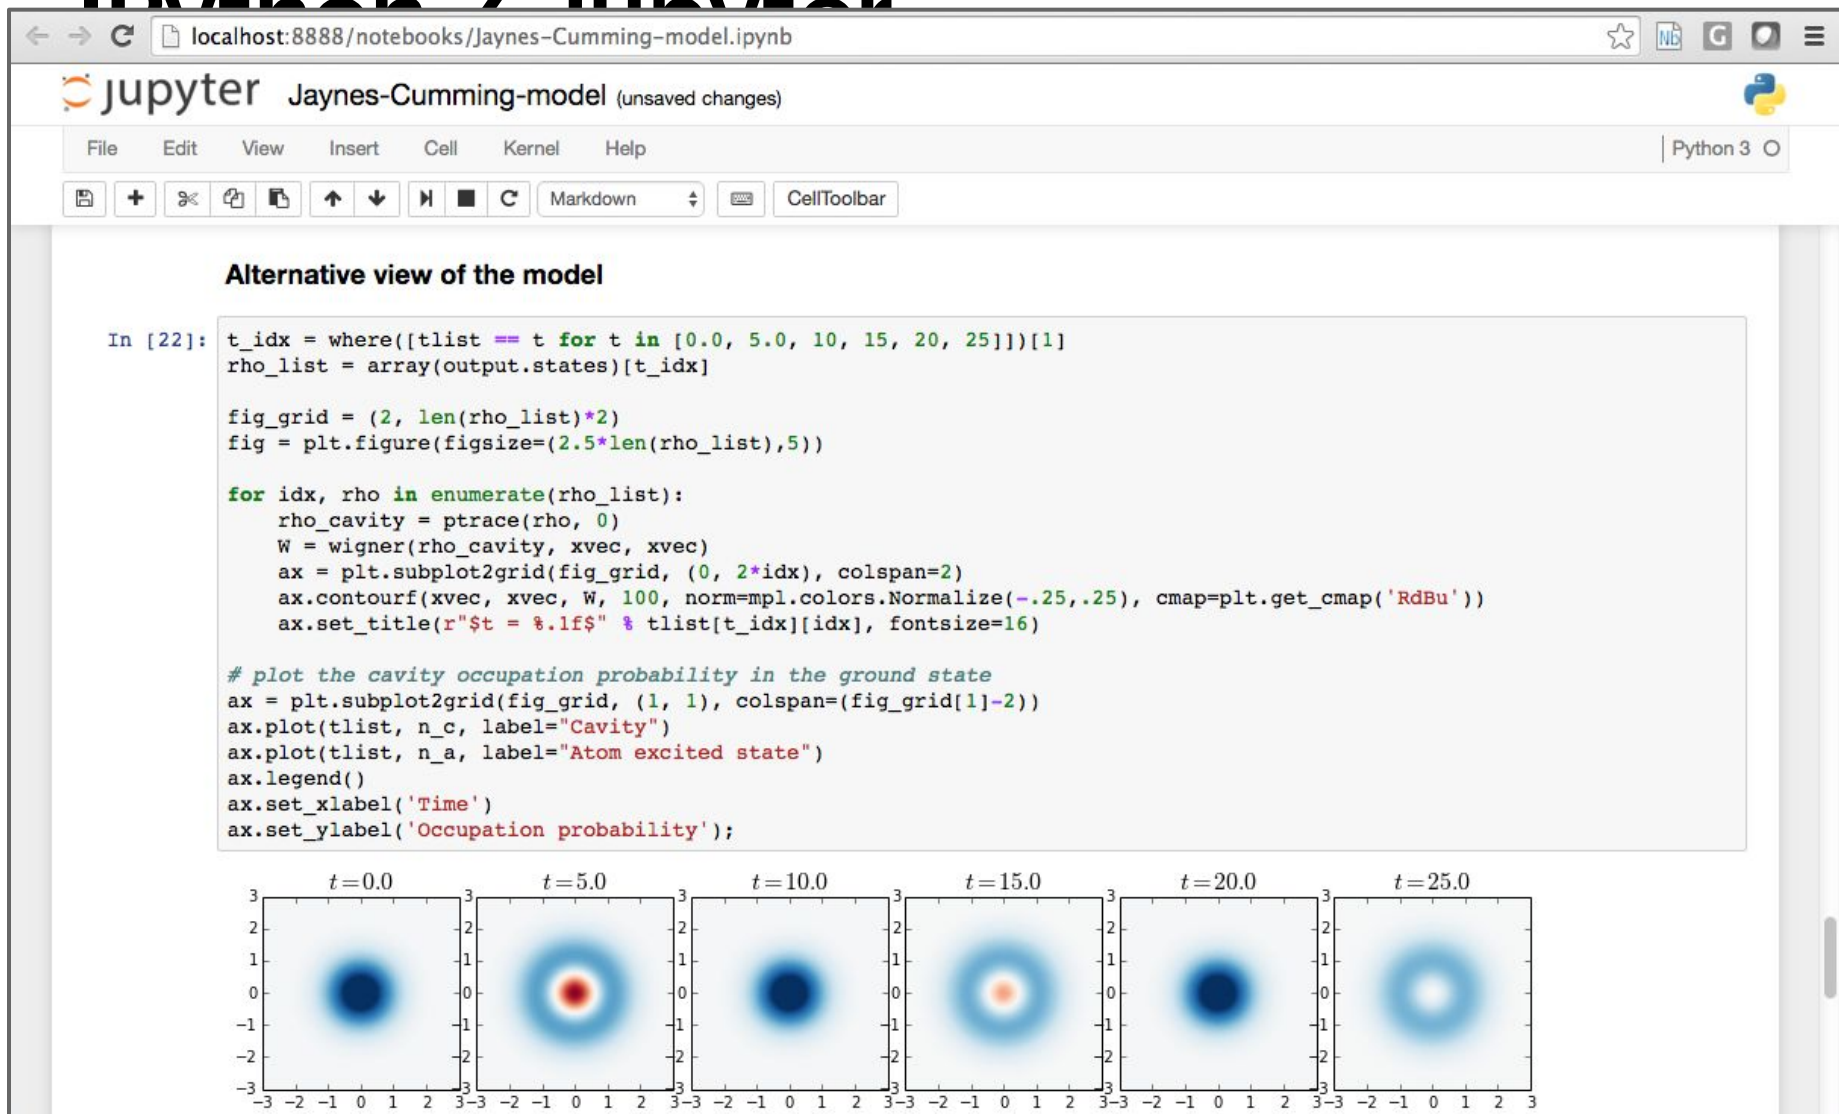


IPython / Jupyter

Terminal, development environment, Notebooks, and more for efficient use of Python in day-to-day work



IPython / Jupyter



IPython




python

TM



IP[y]:
IPython



NumPy



Cython






pythonTM





pandas
 $y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$



NumPy



SymPy

IP[y]:
IPython



SciPy

Provides an interface to common scientific computing Tasks, including wrappers of many NetLib packages.



SciPy

Providing
Tasks

List from <http://docs.scipy.org/doc/scipy/reference/>

- Special functions (**scipy.special**)
- Integration (**scipy.integrate**)
- Optimization (**scipy.optimize**)
- Interpolation (**scipy.interpolate**)
- Fourier Transforms (**scipy.fftpack**)
- Signal Processing (**scipy.signal**)
- Linear Algebra (**scipy.linalg**)
- Sparse Eigenvalue Problems with ARPACK
- Compressed Sparse Graph Routines (**scipy.sparse.csgraph**)
- Spatial data structures and algorithms (**scipy.spatial**)
- Statistics (**scipy.stats**)
- Multidimensional image processing (**scipy.ndimage**)
- File IO (**scipy.io**)


IP[y]:
IPython





Sympy

Library for symbolic computation: algebraic operations, differentiation & integration, optimization, etc.



Sympy

Library for symbolic computation: algebraic operations, differentiation & integration, optimization, etc.

Polynomials and rational functions

SymPy does not expand brackets automatically. The function `expand` is used for this.

```
In [6]: a=(x+y-z)**6  
a
```

```
Out[6]: (x + y - z)6
```


```
In [7]: a=expand(a)  
a
```

```
Out[7]: x6 + 6x5y - 6x5z + 15x4y2 - 30x4yz + 15x4z2 + 20x3y3 - 60x3y2z + 60x3yz2 - 20x3z3  
        + 15x2y4 - 60x2y3z + 90x2y2z2 - 60x2yz3 + 15x2z4 + 6xy5 - 30xy4z + 60xy3z2  
        - 60xy2z3 + 30xyz4 - 6xz5 + y6 - 6y5z + 15y4z2 - 20y3z3 + 15y2z4 - 6yz5 + z6
```



pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$



NumPy



ython



SymPy

IP[y]:
IPython




python™





matplotlib

Matlab-inspired plotting and visualization



matplotlib



NumPy



Cython



SymPy

IP[y]:
IPython



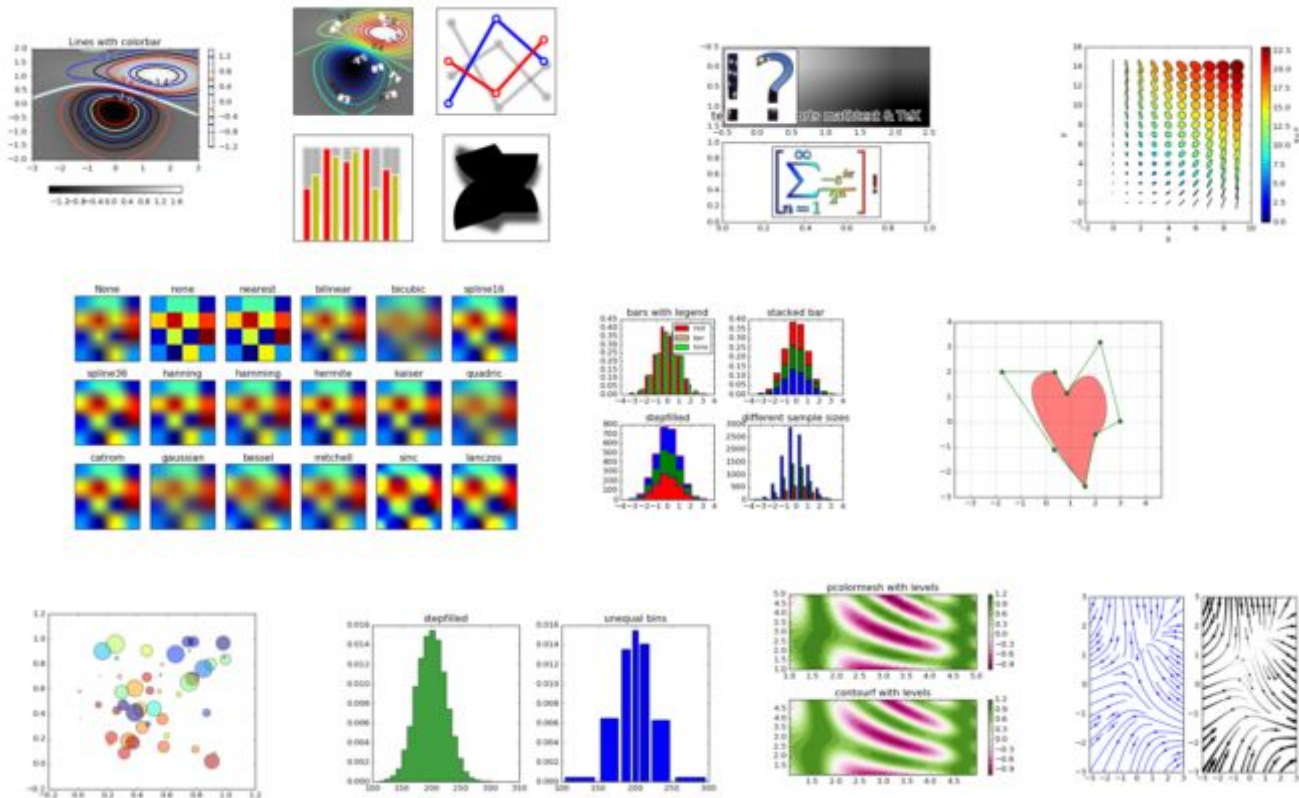
python™



matplotlib



Matlab-inspired plotting and visualization

From <http://matplotlib.org/gallery.html>





pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$


NumPy



ython



SymPy

IP[y]:
IPython




python™



Pandas

R-inspired DataFrames & associated functionality
(data munging & cleaning, group-by & transformations,
and much more)



```
In [1]: import pandas as pd
data = pd.read_csv('iris.csv')
data.head()
```

Out[1]:

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa




```
In [2]: data.groupby('Species').mean()
```

Out[2]:

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
Species				
setosa	5.006	3.428	1.462	0.246
versicolor	5.936	2.770	4.260	1.326
virginica	6.588	2.974	5.552	2.026



pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$


NumPy





SymPy

IP[y]:
IPython








scikit-image
image processing in python



pandas
 $y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$



PyMC



NumPy



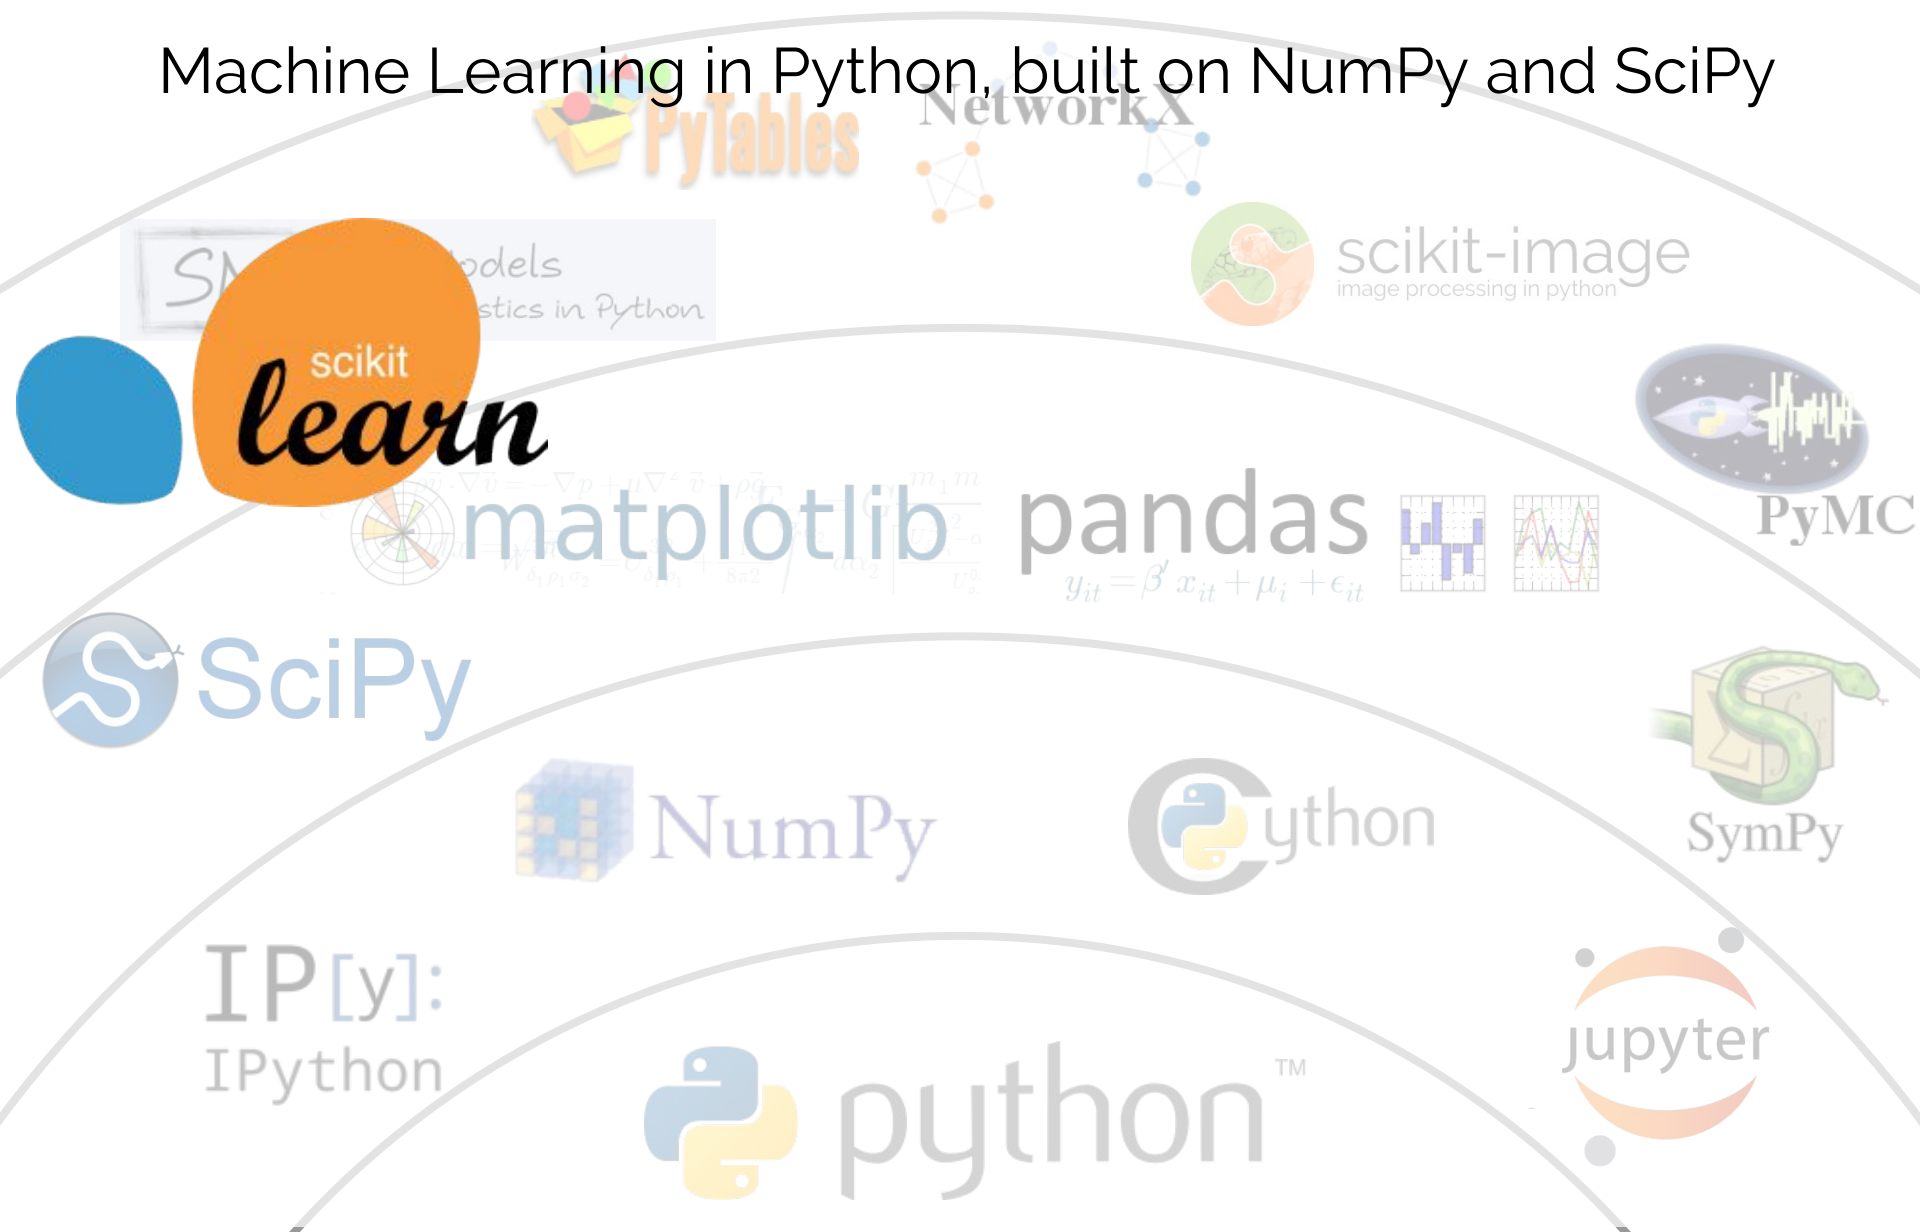
SymPy

IP[y]:
IPython



Scikit-Learn

Machine Learning in Python, built on NumPy and SciPy



Scikit-Learn

Machine Learning in Python, built on NumPy and SciPy

```
In [3]: from sklearn.ensemble import RandomForestClassifier

features = data.drop('Species', axis=1)
labels = data['Species']

model = RandomForestClassifier()
model.fit(features, labels)


predicted = model.predict(features.iloc[:5])
print(predicted)

['setosa' 'setosa' 'setosa' 'setosa' 'setosa']
```

IP[y]:
IPython

 python™

 jupyter








(and many, many more)




scikit-image
image processing in python



pandas
 $y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$



PyMC



NumPy



SymPy

IP[y]:
IPython



Recent-ish Developments

- **Dask**: Parallelization of Data & Computation
- **Numba**: LLVM compilation of Python code
- **Jupyter Lab**: interactive & extensible polyglot development environment
- **Altair**: Declarative Visualization based on Vega-Lite

Dask: Parallel Computation for Distributed Arrays & DataFrames

With minimal changes to your NumPy & Pandas expressions, parallelize your computations over distributed data!

Dask: Parallel Computation for Distributed Arrays & DataFrames

A straightforward NumPy computation:

```
In [3]: import numpy as np

# create an array of normally-distributed random numbers
a = np.random.randn(1000)

# multiply this array by a factor
b = a * 4

# find the minimum value
b_min = b.min()
print(b_min)

-11.4051061336
```

Dask: Parallel Computation for Distributed Arrays & DataFrames

Dask uses the same expressions ...

```
In [4]: import dask.array as da

# create a dask array from the above array
a2 = da.from_array(a, chunks=200)

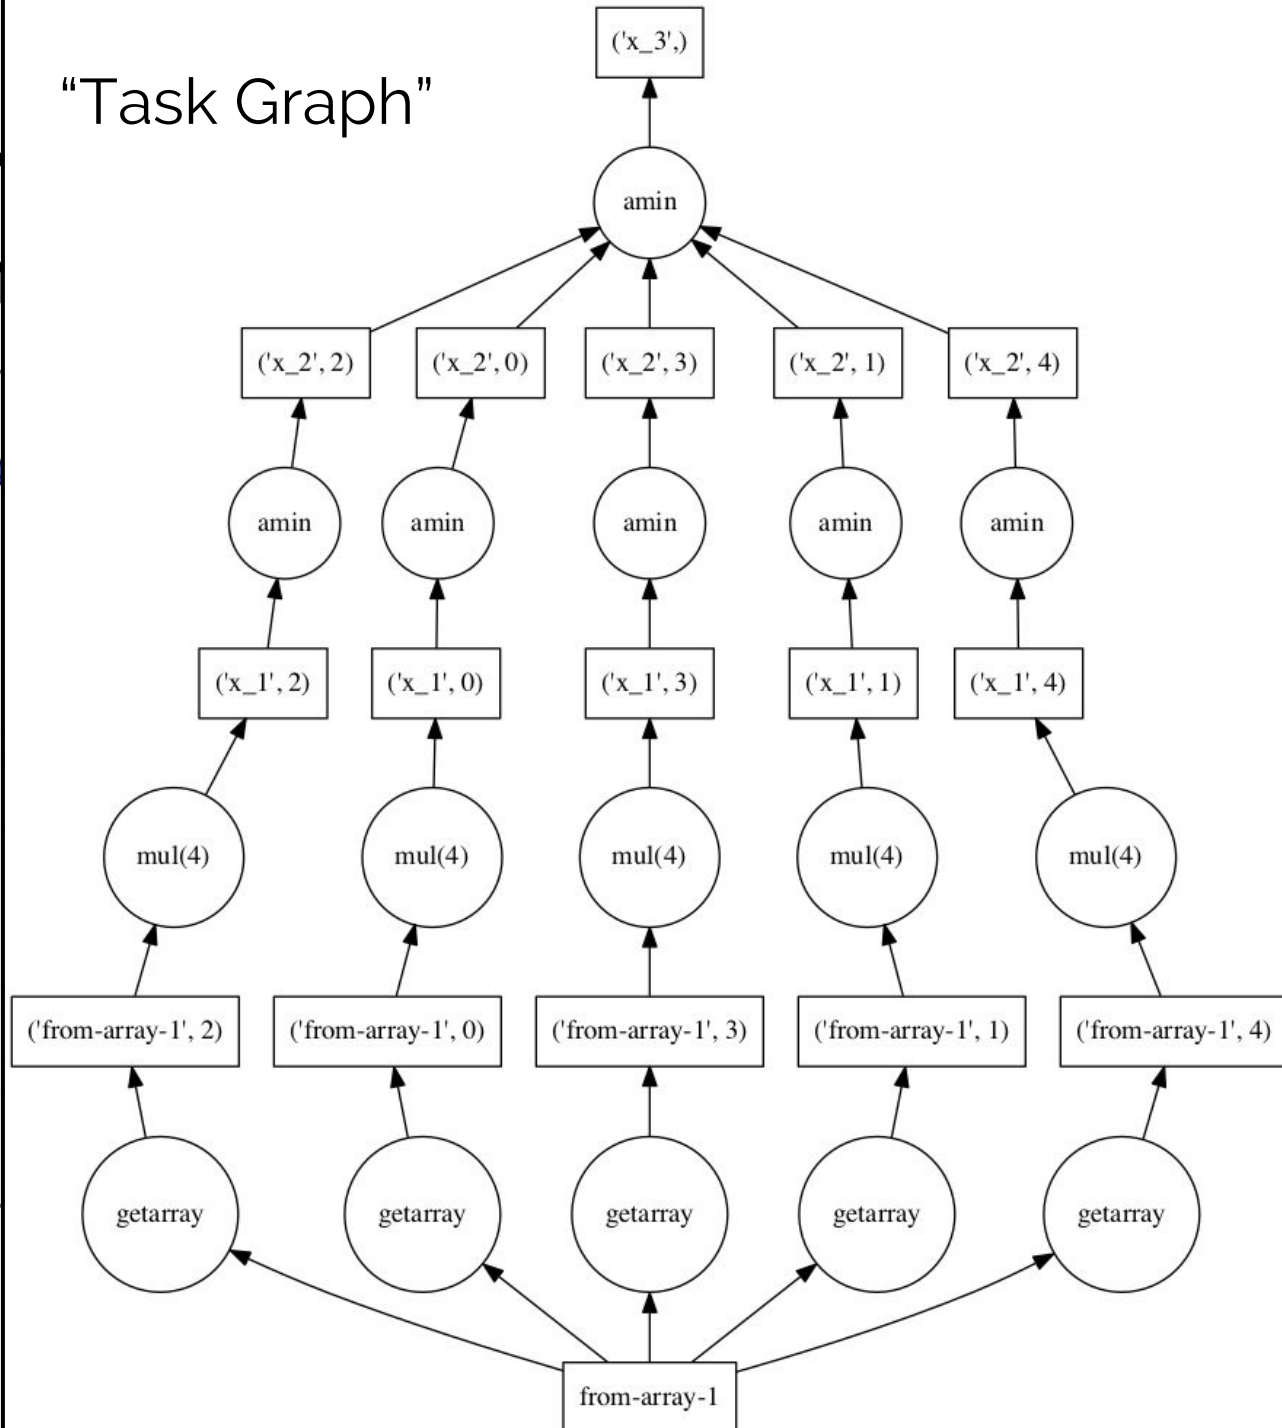
# multiply this array by a factor
b2 = a2 * 4

# find the minimum value
b2_min = b2.min()
print(b2_min)

dask.array<x_3, shape=(), chunks=(), dtype=float64>
```

Task Graph

"Task Graph"



Dask: Parallel Computation for Distributed Arrays & DataFrames

```
In [6]: b2_min.compute()
```

```
Out[6]: -11.405106133564583
```

Numba: JIT-compilation of Python code

With a simple decorator, Python is compiled to LLVM and executes at near C/Fortran speed!

```
def fib(n):  
    a, b = 0, 1  
    for i in range(n):  
        a, b = b, a + b  
    return a
```

```
%timeit fib(50)
```

```
100000 loops, best of 3: 3.83 µs per loop
```

Still some features missing, but very promising (see my blog posts for some examples).

<http://numba.pydata.org/>

Numba: JIT-compilation of Python code

With a simple decorator, Python is compiled to LLVM and executes at near C/Fortran speed!

```
@numba.jit
def fib(n):
    a, b = 0, 1
    for i in range(n):
        a, b = b, a + b
    return a

%timeit(fib(50))
```

1 loops, best of 3: 468 ns per loop

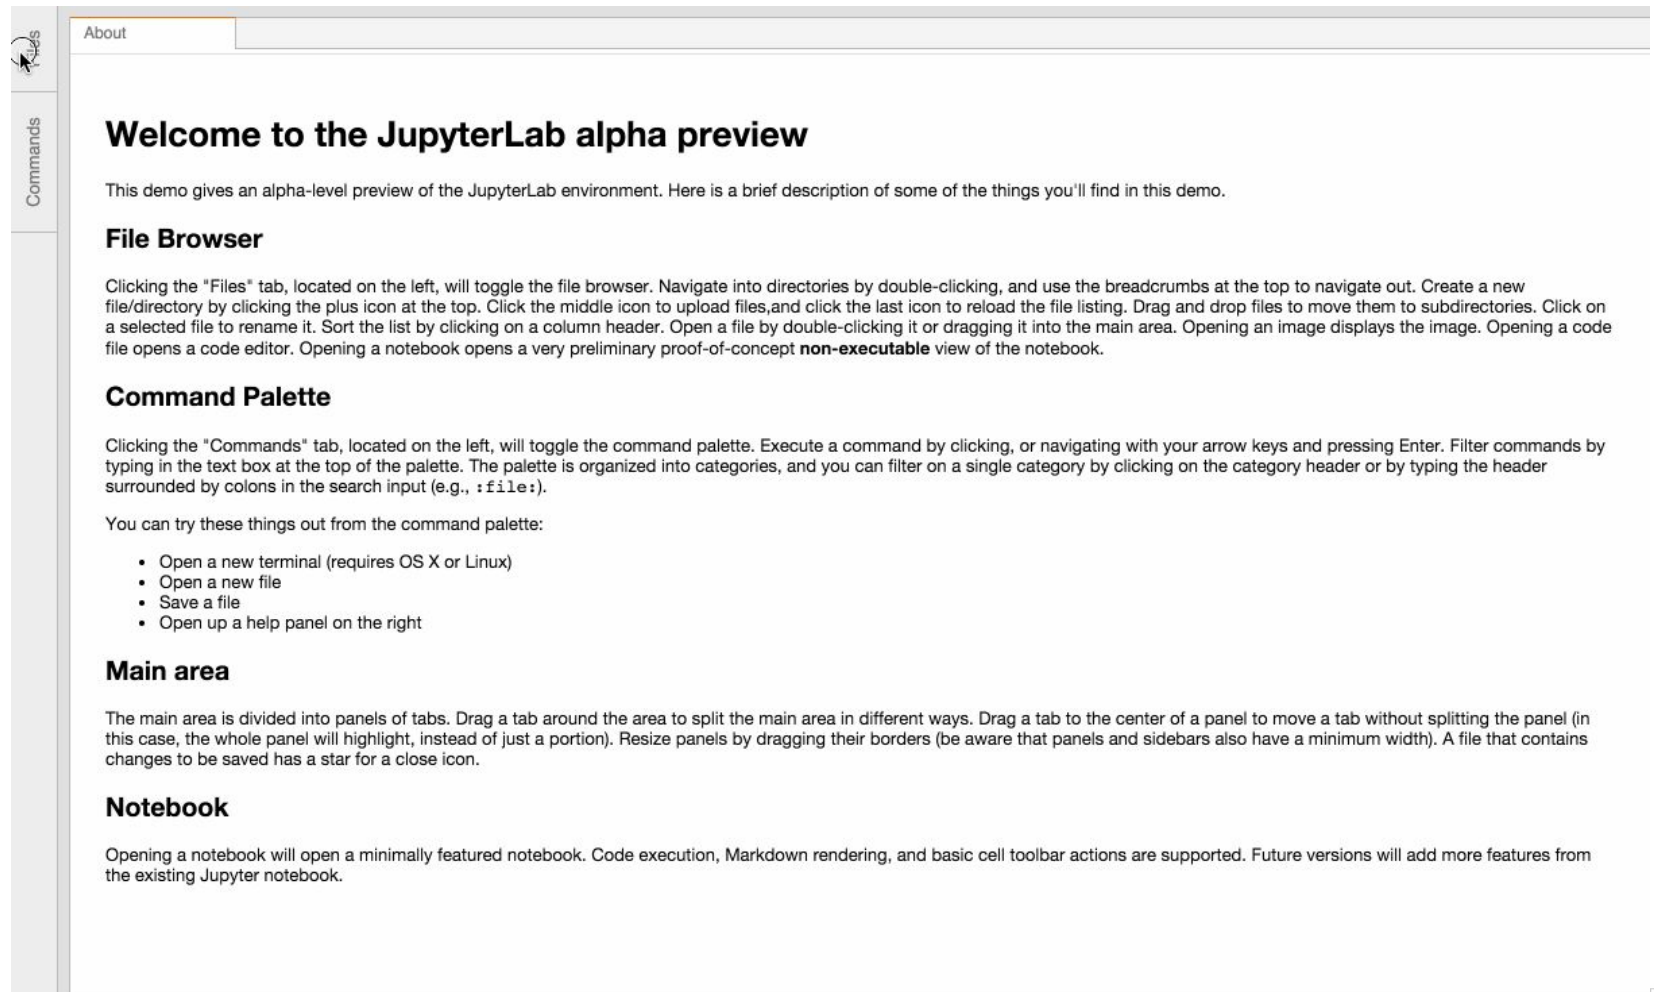
20x speedup!

Still some features missing, but very promising (see my blog posts for some examples).

<http://numba.pydata.org/>

Jupyter Lab

Jupyter beyond notebooks: extensible cross-platform interactive computing environment (release soon!)



Altair: Declarative Visualization based on Vega-Lite

The Visualization story in Python is somewhat confusing . . .

- Matplotlib
- Bokeh
- Plotly
- Seaborn
- Holoviews
- VisPy
- ggplot
- pandas plot
- Lightning

Each library has strengths, but arguably none is yet the “killer viz app” for Data Science.

Most Useful for Data Science is *Declarative Visualization*

Imperative

- Specify *How* something should be done.
- Must manually specify plotting steps
- Specification & Execution intertwined.

Declarative

- Specify *What* should be done
- Details determined automatically
- Separates Specification from Execution

Declarative visualization lets you think about **data** and **relationships**, rather than incidental details.

Enter Altair.

Declarative statistical visualization library for Python,
driven by Vega-Lite

<http://github.com/ellisonbg/altair>

Collaboration with Brian Granger (Jupyter team), myself,
and University of Washington's Interactive Data Lab



UNIVERSITY of WASHINGTON
eScience Institute



#JSM2016



Example: Cars Dataset

```
In [1]: from altair import Chart, load_dataset  
data = load_dataset('cars')
```

```
In [2]: data.head()
```

Out[2]:


	Acceleration	Cylinders	Displacement	Horsepower	Miles_per_Gallon	Name	Origin
0	12.0	8	307.0	130.0	18.0	chevrolet chevelle malibu	USA
1	11.5	8	350.0	165.0	15.0	buick skylark 320	USA
2	11.0	8	318.0	150.0	18.0	plymouth satellite	USA
3	12.0	8	304.0	150.0	16.0	amc rebel sst	USA
4	10.5	8	302.0	140.0	17.0	ford torino	USA

Matplotlib is an *imperative* API:

```
In [4]: import matplotlib.pyplot as plt

def scatter(group):
    plt.plot(group['Horsepower'],
              group['Miles_per_Gallon'],
              'o', label=group.name)


data.groupby('Origin').apply(scatter)
plt.legend(title='Origin')
plt.xlabel('Horsepower')
plt.ylabel('Miles_per_Gallon');
```



Altair is a *declarative* API:

```
In [5]: from altair import datasets, Chart

Chart(data).mark_circle().encode(
    x='Horsepower',
    y='Miles_per_Gallon',
    color='Origin',
)
```



~ ♪

UNIVERSITY of WASHINGTON

Altair is a *declarative* API:

Altair itself contains no renderers, but simply outputs a Vega-Lite visualization specification:

```
from altair.vegalite.v4 import Chart, FacetedChart, Chart

Chart(data).mark_circle().encode(
    x='Horsepower',
    y='Miles_per_Gallon',
    color='Origin',
)
```

```
In [11]: chart.to_dict(data=False)
```

```
Out[11]: {'encoding': {'color': {'field': 'Origin', 'type': 'nominal'},
                        'x': {'field': 'Horsepower', 'type': 'quantitative'},
                        'y': {'field': 'Miles_per_Gallon', 'type': 'quantitative'}},
          'mark': 'circle'}
```

- Portable JSON serialization (Vega-Lite spec)
- Interest from other viz libraries (matplotlib, Bokeh, Plotly) in supporting this serialization.
- Potential for cross-language compatibility



~ ♪


UNIVERSITY of WASHINGTON

Vega-Lite schema is well-defined; allows round-trip between spec and code:

```
In [15]: code = Chart.from_dict(spec).to_altair(data='data')  
print(code)
```


```
Chart(data).mark_circle().encode(  
    color='Origin:N',  
    x='Horsepower:Q',  
    y='Miles_per_Gallon:Q',  
)
```

```
In [16]: eval(code)
```




Altair/Vega-Lite supports many plot types:

```
In [4]: Chart(cars).mark_point().encode(  
    x='Weight_in_lbs',  
    y='Acceleration',  
    color='Cylinders:N'  
)
```




Altair/Vega-Lite supports many plot types:

```
In [6]: Chart(cars).mark_point().encode(  
    x='Displacement',  
    y='Weight_in_lbs',  
    color='Horsepower',  
    row='Origin'  
)  
.configure_cell(width=300, height=150)
```




Altair/Vega-Lite supports many plot types:

```
In [7]: Chart(cars).mark_bar(stacked='normalize').encode(  
    Y('Origin'),  
    X('*:Q', aggregate='count', sort='descending'),  
    Color('Cylinders:N')  
)
```



Altair/Vega-Lite supports many plot types:

```
In [8]: Chart(cars).mark_bar().encode(  
    X('Miles_per_Gallon', bin=Bin(maxbins=20)),  
    Y('*:Q', aggregate='count'),  
    Color('Origin')  
).configure_cell(height=150)
```




Altair/Vega-Lite supports many plot types:

```
In [9]: Chart(cars).mark_text(applyColorToBackground=True).encode(  
    Row('Origin:O'),  
    Column('Cylinders:O'),  
    Color('mean(Miles_per_Gallon):Q', sort='descending'),  
    Text('mean(Miles_per_Gallon):Q')  
)
```



Altair/Vega-Lite supports many plot types:

```
In [10]: Chart(cars).mark_line().encode(  
    X('Year:T', timeUnit='year'),  
    Y('Miles_per_Gallon:Q', aggregate='mean'),  
    Color('Origin:N')  
).configure_cell(height=150)
```



Try Altair:

```
$ conda install altair --channel conda-forge
```

or

```
$ pip install altair  
$ jupyter nbextension install --sys-prefix --py vega
```

For a Jupyter notebook tutorial, type

```
import altair  
altair.tutorial()
```

<http://github.com/ellisonbg/altair/>

Thank You!



Email: jakevdp@uw.edu



Twitter: [@jakevdp](https://twitter.com/jakevdp)



Github: [jakevdp](https://github.com/jakevdp)



Web: <http://vanderplas.com>



Blog: <http://jakevdp.github.io>

Bar Chart: d3

```

var margin = {top: 20, right: 20, bottom: 30, left: 40},
    width = 960 - margin.left - margin.right,
    height = 500 - margin.top - margin.bottom;

var x = d3.scale.ordinal()
    .rangeRoundBands([0, width], .1);

var y = d3.scale.linear()
    .range([height, 0]);

var xAxis = d3.svg.axis()
    .scale(x)
    .orient("bottom");

var yAxis = d3.svg.axis()
    .scale(y)
    .orient("left")
    .ticks(10, "%");

var svg = d3.select("body").append("svg")
    .attr("width", width + margin.left + margin.right)
    .attr("height", height + margin.top + margin.bottom)
    .append("g")
    .attr("transform", "translate(" + margin.left + "," + margin.top + ")");

d3.tsv("data.tsv", type, function(error, data) {
    if (error) throw error;

    x.domain(data.map(function(d) { return d.letter; }));
    y.domain([0, d3.max(data, function(d) { return d.frequency; })]);


    svg.append("g")
        .attr("class", "x axis")
        .attr("transform", "translate(0," + height + ")")
        .call(xAxis);

    svg.append("g")
        .attr("class", "y axis")
        .call(yAxis)
        .append("text")
        .attr("transform", "rotate(-90)")
        .attr("y", 6)
        .attr("dy", ".71em")
        .style("text-anchor", "end")
        .text("Frequency");

    svg.selectAll(".bar")
        .data(data)
        .enter().append("rect")
        .attr("class", "bar")
        .attr("x", function(d) { return x(d.letter); })
        .attr("width", x.rangeBand())
        .attr("y", function(d) { return y(d.frequency); })
        .attr("height", function(d) { return height - y(d.frequency); });
});


function type(d) {
    d.frequency = +d.frequency;
    return d;
}

```



Bar Chart: Vega

```
{
  "width": 400,
  "height": 200,
  "padding": {"top": 10, "left": 30, "bottom": 30, "right": 10},
  "data": [
    {
      "name": "table",
      "values": [
        {"x": 1, "y": 28}, {"x": 2, "y": 55},
        {"x": 3, "y": 43}, {"x": 4, "y": 91},
        {"x": 5, "y": 81}, {"x": 6, "y": 53},
        {"x": 7, "y": 19}, {"x": 8, "y": 87},
        {"x": 9, "y": 52}, {"x": 10, "y": 48},
        {"x": 11, "y": 24}, {"x": 12, "y": 49},
        {"x": 13, "y": 87}, {"x": 14, "y": 66},
        {"x": 15, "y": 17}, {"x": 16, "y": 27},
        {"x": 17, "y": 68}, {"x": 18, "y": 16},
        {"x": 19, "y": 49}, {"x": 20, "y": 15}
      ]
    }
  ],
  "scales": [
    {
      "name": "x",
      "type": "ordinal",
      "range": "width",
      "domain": {"data": "table", "field": "x"}
    },
    {
      "name": "y",
      "type": "linear",
      "range": "height",
      "domain": {"data": "table", "field": "y"},
      "nice": true
    }
  ],
  "axes": [
    {"type": "x", "scale": "x"},
    {"type": "y", "scale": "y"}
  ],
  "marks": [
    {
      "type": "rect",
      "from": {"data": "table"},
      "properties": {
        "enter": {
          "x": {"scale": "x", "field": "x"},
          "width": {"scale": "x", "band": true, "offset": -1},
          "y": {"scale": "y", "field": "y"},
          "y2": {"scale": "y", "value": 0}
        }
      },
      "update": {
        "fill": {"value": "steelblue"}
      }
    }
  ]
}
```



```

{
  "description": "A simple bar chart with embedded data.",
  "data": {
    "values": [
      {"a": "A", "b": 28}, {"a": "B", "b": 55}, {"a": "C", "b": 43},
      {"a": "D", "b": 91}, {"a": "E", "b": 81}, {"a": "F", "b": 53},
      {"a": "G", "b": 19}, {"a": "H", "b": 87}, {"a": "I", "b": 52}
    ]
  },
  "mark": "bar",
  "encoding": {
    "x": {"field": "a", "type": "ordinal"},
    "y": {"field": "b", "type": "quantitative"}
  }
}

```

Bar Chart: Vega-Lite

